

Compressive Sensing in Video Encoding

Brian Azizi


Laboratory for Scientific Computing, University of Cambridge

May 18, 2016

Thanks to: Dr Anita Faul (Supervisor)
Dr Nikos Nikiforakis
Georgios Pilikos

- Project Outcome & Demonstration
- Theoretical Background
- Further Work
- More demos of results
- Bibliography

2D Demonstration



Corrupted Image (70% missing)



Scale 1 Reconstruction

See 2014 MPhil Thesis by Georgios Pilikos

2D Demonstration



Cascade to scale 2



Cascade to scale 3

See 2014 MPhil Thesis by Georgios Pilikos

Current prototype:



Corrupted Video (60% missing)



Scale 1 Reconstruction

3D Demonstration

- No cascade of reconstructions yet - ask me later this week
- For reference, the original video:



- Three building blocks:
 - *Compressive Sensing* - Novel signal processing framework allowing for near-perfect reconstruction of heavily under-sampled signals
 - *Discrete Wavelet Transforms* - Required for obtaining sparse representations of the signals
 - *Relevance Vector Machine* - Machine Learning Algorithm for performing the reconstruction via regression

Compressive Sensing: Reconstruction

- General idea: Reconstruct signal $\mathbf{x} \in \mathbb{R}^N$ from measurements $\mathbf{y} \in \mathbb{R}^M$ where $M \ll N$
- Corresponds to solving under-determined linear system $\Omega \mathbf{x} = \mathbf{y}$, $\Omega \in \mathbb{R}^{M \times N}$
- In general, shouldn't be possible by fundamental theorem of Linear Algebra (“as many equations as unknowns”)

- However, for a certain class of signals it is possible to get perfect reconstruction. Namely if
 - 1 x is *sparse* (i.e. most elements are zero), or
 - 2 it is possible to change basis so that the transformed signal $w = \Psi^T x$ is sparse
- So in the later case: $y = \Omega x = \Omega \Psi w \equiv \Phi w$
- We know $y \in \mathbb{R}^M$ and $\Phi \in \mathbb{R}^{M \times N}$ and want $w \in \mathbb{R}^N$

Compressive Sensing: Deterministic Methods

- Ideally, we want $\hat{\mathbf{w}} = \arg \min ||\mathbf{w}'||_0$ such that $\Phi \mathbf{w}' = \mathbf{y}$
- $||\mathbf{v}||_0$ is the L_0 norm = number of non-zeros entries of \mathbf{v}
- This would give the sparsest solution

Compressive Sensing: Deterministic Methods

- Ideally, we want $\hat{\mathbf{w}} = \arg \min ||\mathbf{w}'||_0$ such that $\Phi \mathbf{w}' = \mathbf{y}$
- $||\mathbf{v}||_0$ is the L_0 norm = number of non-zeros entries of \mathbf{v}
- This would give the sparsest solution
- But it turns out to be NP-complete - hence not feasible

Compressive Sensing: Deterministic Methods


- Ideally, we want $\hat{\mathbf{w}} = \arg \min ||\mathbf{w}'||_0$ such that $\Phi \mathbf{w}' = \mathbf{y}$
- $||\mathbf{v}||_0$ is the L_0 norm = number of non-zeros entries of \mathbf{v}
- This would give the sparsest solution
- But it turns out to be NP-complete - hence not feasible
- Recently[1], there's been much success by doing L_1 instead of L_0 minimization: $\hat{\mathbf{w}} = \arg \min ||\mathbf{w}'||_1$ such that $\Phi \mathbf{w}' = \mathbf{y}$
- $||\mathbf{v}||_1 = \sum_{i=1}^N |v_i|$ is the L_1 norm

Compressive Sensing: Deterministic Methods

- Ideally, we want $\hat{\mathbf{w}} = \arg \min \|\mathbf{w}'\|_0$ such that $\Phi \mathbf{w}' = \mathbf{y}$
- $\|\mathbf{v}\|_0$ is the L_0 norm = number of non-zeros entries of \mathbf{v}
- This would give the sparsest solution
- But it turns out to be NP-complete - hence not feasible
- Recently[1], there's been much success by doing L_1 instead of L_0 minimization: $\hat{\mathbf{w}} = \arg \min \|\mathbf{w}'\|_1$ such that $\Phi \mathbf{w}' = \mathbf{y}$
- $\|\mathbf{v}\|_1 = \sum_{i=1}^N |v_i|$ is the L_1 norm
- These deterministic approaches are not relevant for us at the moment
- we will use the RVM

Compressive Sensing: Acquisition

- Beside reconstruction, a large part of Compressive Sensing is concerned with efficient *acquisition* of signals, i.e. how to measure y directly without measuring the entirety of x first
- Not relevant for us at this stage
- Instead, we simulate y by taking x and randomly deleting a certain percentage of its entries



- E.g. for 2D signals: $x =$ and $y =$
- Then attempt to reconstruct x from y

- Technical Aside: I talk somewhat interchangeably about 2D signals (e.g. images) and 3D signals (e.g. videos)
- Under the hood, we actually need to store our signals as 1-dimensional vectors
- Reason: RVM operates on vectors
- For images, we stack the columns on top of each other to form one long vector in \mathbb{R}^{hw} (height times width)
- For video, we stack the columns in each frame and then stack the frames ($\Rightarrow \mathbf{x} \in \mathbb{R}^{hwf}$)

Sparse Representations


- It is possible to obtain sparse representations for a large variety of signals
- Common example in signal processing: Transforming sounds from time domain to frequency domain via (Discrete) Fourier Transform
- For natural images and image sequences (videos), there are two widely used transforms:
 - Discrete Haar Wavelet Transform
 - Discrete Cosine Transform
- There exist more transforms and probably better ones
- Ideally, we would like the data to tell us which one to use via some kind of Deep Learning - but this is getting off-topic

2D Haar Wavelets

- Current prototype uses Haar wavelets at the first scale
- Very simple: transform essentially consists of taking averages and differences of patches of pixels
- At scale 1, the individual basis functions have support over 2×2 patches

2D Haar Wavelets

- 2D Haar wavelets at scale 1 can be visualized by




- White $\equiv +\frac{1}{2}$, black $\equiv -\frac{1}{2}$, dark gray $\equiv 0$
- So for example $x_{2,7} \rightarrow w_{2,7} = \frac{1}{2}(x_{3,5} - x_{3,6} + x_{4,5} - x_{4,6})$

2D Haar Wavelets: Example

- Haar wavelet transform of




is $w =$



- Most entries in w are zero or very close to zero (shown as black)
 $\Rightarrow w$ is sparse

2D Haar Wavelets: Example

- Somewhat easier to see what's going on if we invert and translate the colours a little bit:



- Top left is an average of x , remaining entries in w are detail coefficients that capture horizontal, vertical and diagonal *edges* in x

3D Haar Wavelets

- Think of a video as a volume created by stacking the individual frames
- To take the Haar wavelet transform of a video x we have two approaches:
 - Take the 2D transform of each frame individually \rightarrow process x as a sequence of independent images
 - Use 3D Haar wavelets and process video as a volume \rightarrow exploits continuity between frames
- We used the 3D wavelets approach because it is more general
- 3D Haar wavelets have support over $2^j \times 2^j \times 2^j$ blocks (at the j th scale)

3D Haar Wavelets

Demo of Haar wavelet transform of soccer.yuv

$x =$



(30 frames per second)

$w =$



(20 frames per second)



- Recall the goal of Compressive Sensing: Find sparsest solution to $\Phi \mathbf{w} = \mathbf{y}$
- We mentioned deterministic methods (L_1 minimization)
- But we will use a probabilistic method by treating it as a Machine Learning problem
- Specifically, we use the RVM to do a regression:
 - 1 Input a *target vector* $\mathbf{y} \in \mathbb{R}^M$
 - 2 Input a *design matrix* $\Phi \in \mathbb{R}^{M \times N}$
 - 3 The RVM outputs a *sparse coefficients vector* $\mathbf{w}^* \in \mathbb{R}^N$

- More specifically, RVM gives us a *posterior distribution* for \mathbf{w}^* :
 $\mathbf{w}^* | \text{data} \sim \mathcal{N}(\boldsymbol{\mu}, \Sigma)$
- The *posterior mean* $\boldsymbol{\mu}$ is usually very sparse
- Reconstruct the original signal as $\mathbf{x}^* = \Psi \boldsymbol{\mu}$
- Even better: we can use Σ to get error bars in our reconstruction
- We can use error bars to create a cascade of reconstructions and boost the quality
- Focus for next couple of days

The Relevance Vector Machine: Theory

- For our purposes, it is okay to treat the RVM as a “black box”
- However, if interested in theory:
- The RVM [2] is a Bayesian Machine Learning algorithm
- It was developed by Mike Tipping[2] and later improved upon by him and Anita Faul[3]

- It models the data as $p(y_i | \mathbf{x}, \Phi) = \mathcal{N}(y_i | \mathbf{w}^T \boldsymbol{\phi}(x_i), \sigma^2)$
- Bayesian, so put a prior on \mathbf{w} : $p(\mathbf{w}) = \prod_{j=1}^N \mathcal{N}(w_j | 0, \alpha_j^{-1})$
- Use training data (observations) to obtain a posterior for \mathbf{w} :
 $p(\mathbf{w}) = \mathcal{N}(\mathbf{w} | \boldsymbol{\mu}, \Sigma)$
- During training, many α_j become infinite $\Rightarrow w_j = 0$ with infinite precision
- Thus \mathbf{w} is *sparse*

- Progress so far:
 - Implemented a working prototype
 - Use Haar Wavelet transform
 - Got some initial results
- Still to do:
 - Multi-scale Cascade of RVMs
 - Try more Basis Functions, in particular the Discrete Cosine Transform
 - Study: When does it work well, when not?
 - Lit review: How does it compare to other existing methods?
 - How can it be improved?

Questions?

More demos

Flickering lines (60% missing data)

$y =$



$x^* =$



Uniform noise (60% missing data)




$y =$



$x^* =$



Selected References

-  E. J. Candes, T. Tao, Decoding by linear programming, Information Theory, IEEE Transactions on 51 (12) (2005) 4203–4215.
-  M. E. Tipping, Sparse bayesian learning and the relevance vector machine, The journal of machine learning research 1 (2001) 211–244.
-  M. E. Tipping, A. C. Faul, et al., Fast marginal likelihood maximisation for sparse bayesian models., in: AISTATS, 2003.